

Ex 3

**Implement a simple map-reduce code for the wordcount problem using Java/Python.
(Create the jar files and run the code using HDFS.)**

```
su - hadoop
```

```
mkdir ~/WordCountProject  
cd ~/WordCountProject
```

```
nano WordCountMapper.java
```

```
import java.io.IOException;  
import org.apache.hadoop.io.IntWritable;  
import org.apache.hadoop.io.LongWritable;  
import org.apache.hadoop.io.Text;  
import org.apache.hadoop.mapreduce.Mapper;  
  
public class WordCountMapper extends Mapper<LongWritable, Text, Text, IntWritable> {  
    private final static IntWritable one = new IntWritable(1);  
    private Text word = new Text();  
  
    @Override  
    public void map(LongWritable key, Text value, Context context) throws IOException,  
        InterruptedException {  
        String[] words = value.toString().split("\\s+");  
        for (String str : words) {  
            word.set(str);  
            context.write(word, one);  
        }  
    }  
}
```

```
nano WordCountReducer.java
```

```
import java.io.IOException;  
import org.apache.hadoop.io.IntWritable;  
import org.apache.hadoop.io.Text;  
import org.apache.hadoop.mapreduce.Reducer;
```

```

public class WordCountReducer extends Reducer<Text, IntWritable, Text, IntWritable> {

    @Override
    public void reduce(Text key, Iterable<IntWritable> values, Context context) throws
IOException, InterruptedException {
        int sum = 0;
        for (IntWritable val : values) {
            sum += val.get();
        }
        context.write(key, new IntWritable(sum));
    }
}

```

nano WordCount.java

```

import org.apache.hadoop.conf.Configuration;
import org.apache.hadoop.fs.Path;
import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapreduce.Job;
import org.apache.hadoop.mapreduce.lib.input.FileInputFormat;
import org.apache.hadoop.mapreduce.lib.output.FileOutputFormat;

public class WordCount {

    public static void main(String[] args) throws Exception {
        Configuration conf = new Configuration();
        Job job = Job.getInstance(conf, "word count");

        job.setJarByClass(WordCount.class);
        job.setMapperClass(WordCountMapper.class);
        job.setCombinerClass(WordCountReducer.class);
        job.setReducerClass(WordCountReducer.class);

        job.setOutputKeyClass(Text.class);
        job.setOutputValueClass(IntWritable.class);

        FileInputFormat.addInputPath(job, new Path(args[0]));
        FileOutputFormat.setOutputPath(job, new Path(args[1]));

        System.exit(job.waitForCompletion(true) ? 0 : 1);
    }
}

```

```
mkdir wordcount_classes
```

```
javac -classpath $(hadoop classpath) -d wordcount_classes WordCountMapper.java  
WordCountReducer.java WordCount.java
```

```
jar -cvf WordCount.jar -C wordcount_classes/ .
```

```
start-dfs.sh
```

```
start-yarn.sh
```

```
echo "Hello Hadoop Hello MapReduce" > input.txt
```

```
whoami (find username)
```

```
(hadoop - username)
```

```
hdfs dfs -mkdir /user/hadoop/input
```

```
hdfs dfs -put input.txt /user/hadoop/input
```

```
hadoop jar WordCount.jar WordCount /user/hadoop/input /user/hadoop/output
```

```
hdfs dfs -cat /user/hadoop/output/part-r-00000
```

Delete existing dic -error

```
hdfs dfs -rm -r /user/hadoop/output
```

```
hadoop jar WordCount.jar WordCount /user/hadoop/input /user/hadoop/output
```

```
hdfs dfs -cat /user/hadoop/output/part-r-00000
```

Python version

```
mkdir ~/WordCountPythonProject
```

```
nano mapper.py
```

```
# mapper.py
import sys

# Input comes from standard input (line by line)
for line in sys.stdin:
    line = line.strip() # Remove leading and trailing whitespace
    words = line.split() # Split line into words

    # Output each word with a count of 1
    for word in words:
        print(f'{word}\t1')
```

nano reducer.py

```
# reducer.py
import sys

current_word = None
current_count = 0
word = None

# Input comes from standard input
for line in sys.stdin:
    line = line.strip()
    word, count = line.split('\t', 1)

    try:
        count = int(count)
    except ValueError:
        continue

    # Sum counts for each word
    if current_word == word:
        current_count += count
    else:
        if current_word:
            print(f'{current_word}\t{current_count}')
            current_word = word
            current_count = count

# Output the last word
if current_word == word:
    print(f'{current_word}\t{current_count}')
```

```
start-dfs.sh
start-yarn.sh
```

(Delete file output)

```
hdfs dfs -rm /user/hadoop/input/input.txt
```

```
echo "Hello Hadoop Hello MapReduce" > input1.txt
```

```
hdfs dfs -put input1.txt /user/hadoop/input
```

```
hadoop jar $HADOOP_HOME/share/hadoop/tools/lib/hadoop-streaming-*.jar \
-input /user/hadoop/input \
-output /user/hadoop/output \
-mapper "python3 mapper.py" \
-reducer "python3 reducer.py"
```

```
hdfs dfs -cat /user/hadoop/output/part-00000
```

Ex 4

1. Implement map reduce for NCDC weather dataset using Hadoop -fine the max and min temperature.

```
start-dfs.sh
start-yarn.sh
```

2. Implement Apriori algorithm using map reduce paradigm.

```
start-dfs.sh
start-yarn.sh
```

```
# Create a directory for your input data in HDFS
hadoop fs -mkdir -p /user/hadoop/input
```

```
# Create the transactions file
```

```
echo -e "milk,bread,butter\nbread,butter,juice\nmilk,juice\nbread,milk,juice" > transactions.txt
```

```
# Upload the input file to HDFS
```

```
hadoop fs -put transactions.txt /user/hadoop/input
```

```
mkdir AprioriMR
```

```
cd AprioriMR
```

```
mkdir src
```

```
cd src
```

AprioriMR.java (create inside src)

```
import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapreduce.Mapper;
import org.apache.hadoop.mapreduce.Reducer;
import org.apache.hadoop.conf.Configuration;
import org.apache.hadoop.fs.Path;
import org.apache.hadoop.mapreduce.Job;
import org.apache.hadoop.mapreduce.lib.input.FileInputFormat;
import org.apache.hadoop.mapreduce.lib.output.FileOutputFormat;
```

```
import java.io.IOException;
import java.util.StringTokenizer;
```

```
public class AprioriMR {
```

```
    public static class ItemsetMapper extends Mapper<Object, Text, Text, IntWritable> {
        private final static IntWritable one = new IntWritable(1);
        private Text item = new Text();
```

```
        public void map(Object key, Text value, Context context) throws IOException,
        InterruptedException {
            StringTokenizer itr = new StringTokenizer(value.toString(), ",");
            while (itr.hasMoreTokens()) {
                item.set(itr.nextToken().trim());
                context.write(item, one);
            }
        }
    }
}
```

```
    public static class ItemsetReducer extends Reducer<Text, IntWritable, Text, IntWritable> {
        private int minSupport;
```

```

@Override
protected void setup(Context context) {
    Configuration conf = context.getConfiguration();
    minSupport = conf.getInt("minSupport", 2); // Example threshold
}

public void reduce(Text key, Iterable<IntWritable> values, Context context) throws
IOException, InterruptedException {
    int sum = 0;
    for (IntWritable val : values) {
        sum += val.get();
    }
    if (sum >= minSupport) {
        context.write(key, new IntWritable(sum));
    }
}
}

public static void main(String[] args) throws Exception {
    Configuration conf = new Configuration();
    conf.setInt("minSupport", 2); // Set minimum support here
    Job job = Job.getInstance(conf, "apriori");
    job.setJarByClass(AprioriMR.class);
    job.setMapperClass(ItemsetMapper.class);
    job.setCombinerClass(ItemsetReducer.class);
    job.setReducerClass(ItemsetReducer.class);
    job.setOutputKeyClass(Text.class);
    job.setOutputValueClass(IntWritable.class);

    FileInputFormat.addInputPath(job, new Path(args[0]));
    FileOutputFormat.setOutputPath(job, new Path(args[1]));

    System.exit(job.waitForCompletion(true) ? 0 : 1);
}
}

```

```

# Navigate to the source directory
cd src

```

```

# Compile the code
javac -classpath `hadoop classpath` -d ../ AprioriMR.java

```

```

# Go back to the project root
cd ..

```

Package the compiled code into a JAR file

```
jar -cvf AprioriMR.jar -C . .
```

Run the Hadoop job

```
hadoop jar AprioriMR.jar AprioriMR /user/hadoop/input /user/hadoop/output
```

List the output files in HDFS

```
hadoop fs -ls /user/hadoop/output
```

View the results

```
hadoop fs -cat /user/hadoop/output/part-r-00000
```

(Troubleshooting)

Remove input and output directories from HDFS

```
hadoop fs -rm -r /user/hadoop/input
```

```
hadoop fs -rm -r /user/hadoop/output
```

List files in the src directory to check if AprioriMR.java is there

```
ls src
```

Navigate to the src directory

```
cd src
```

Compile the code using Hadoop's classpath

```
javac -classpath `hadoop classpath` -d ../ AprioriMR.java
```

(if already output file is there)

```
hadoop fs -rm -r /user/hadoop/output
```

```
hadoop jar AprioriMR.jar AprioriMR /user/hadoop/input /user/hadoop/output
```

```
File Output Format Counters
  Bytes Written=32
hadoop@CynddiaPC:~/AprioriMR$ # List files in the output directory to verify the job's completion
hadoop fs -ls /user/hadoop/output

# Display the contents of the result file
hadoop fs -cat /user/hadoop/output/part-r-00000
Found 2 items
-rw-r--r--  1 hadoop supergroup          0 2024-11-03 12:00 /user/hadoop/output/_SUCCESS
-rw-r--r--  1 hadoop supergroup        32 2024-11-03 12:00 /user/hadoop/output/part-r-00000
bread      3
butter     2
juice      3
milk       3
hadoop@CynddiaPC:~/AprioriMR$
```


Ex 5

Installing pyspark + jupyter

[How to Run PySpark on Jupyter Notebook | phoenixNAP KB](#)

Pyspark with jupyter

<https://chatgpt.com/share/6727240e-8e94-8013-a25c-b04e1d94de0d>

1. Run the wordcount program that you did using hadoop usingpyspark.

```
import os
import findspark
from pyspark.sql import SparkSession

# Set environment variables
os.environ['SPARK_HOME'] = "/home/hadoop/.local/lib/python3.10/site-packages/pyspark" #
Adjust this if necessary
os.environ['HADOOP_HOME'] = "/path/to/hadoop" # If you have Hadoop installed
os.environ['PYSPARK_PYTHON'] = "python3" # Or "python" depending on your setup

# Initialize findspark to find the Spark installation
findspark.init()

# Create a Spark session
spark = SparkSession.builder \
    .appName("WordCount") \
    .getOrCreate()

# Print Spark version to confirm it's working
print(spark.version)

# Example: Word Count
text_data = ["Hello world", "Hello Spark", "Hello Jupyter"]
rdd = spark.sparkContext.parallelize(text_data)
```

```
word_counts = rdd.flatMap(lambda line: line.split(" ")) \
    .map(lambda word: (word, 1)) \
    .reduceByKey(lambda a, b: a + b)
```

```
# Collect and print results
for word, count in word_counts.collect():
    print(f"{word}: {count}")
```

Movielens dataset - find out for each movie, how are the ratings distributed

(create txt movielens dataset in jupyter directory - not local file)

```
from pyspark.sql import SparkSession
from pyspark.sql.functions import col, count
```

```
# Create a Spark session
spark = SparkSession.builder \
    .appName("MovieLens Ratings Distribution") \
    .getOrCreate()
```

```
data_path = "Movielens.txt"
movies_df = spark.read.csv(data_path, sep='\t', inferSchema=True) \
    .toDF("user_id", "movie_id", "rating", "timestamp")
# Calculate the distribution of ratings for each movie
rating_distribution = movies_df.groupBy("movie_id", "rating") \
    .agg(count("rating").alias("rating_count")) \
    .orderBy("movie_id", "rating")
```

```
rating_distribution.show()
```

Ex 6

1. Use the "friends_test" dataset. Col1 is ID, Col2 is name, Col 3 is Age, Col 4 is num of friends. Understand mapvalues function of RDD in spark and find the average number of friends for each unique age present in the dataset.

```
import pyspark
from pyspark.sql import SparkSession
spark =
SparkSession.builder.master('local').appName('friends_dataset').getOrCreate()
sc=spark.sparkContext

# Load dataset
data_path='friends_test.csv'
rdd=sc.textFile(data_path)

age_friends_rdd = rdd.map(lambda line: line.split(',')) \
                      .map(lambda cols: (int(cols[2]), (int(cols[3]), 1)))

# Sum up num_friends and count for each age
age_friends_totals = age_friends_rdd.reduceByKey(lambda a, b: (a[0] +
b[0], a[1] + b[1]))

# Calculate the average number of friends for each age
average_friends_by_age = age_friends_totals.mapValues(lambda total:
total[0] / total[1])

# Collect and display the results
results = average_friends_by_age.collect()
for age, avg_friends in results:
    print(f"Age: {age}, Average Friends: {avg_friends:.2f}")
```

2. Use the "temp.csv" dataset. Column headers are present in the dataset. Understand filter operations and filter out only the "TMIN" values from the "desc" column. With the resultant data (RDD) find the following: a. Minimum temperature (overall) b. Minimum temperature for every ItemID c. Minimum temperature for every StationID.

Use the same dataset, filter only "TMAX" column and find the maximum temperatures just like the ones mentioned above.

```
from pyspark.sql import SparkSession

# Initialize Spark session
spark = SparkSession.builder \
    .appName("Temp_dataset") \
    .getOrCreate()

# Load dataset
data_path = "temp.csv" # Replace with actual path
rdd = spark.sparkContext.textFile(data_path)

# Extract the header
header = rdd.first()

data_rdd = rdd.filter(lambda row: row != header) # Remove the header
# Split each row by comma and convert to (StationID, ItemID, desc, temp)
format
data_rdd = data_rdd.map(lambda line: line.split(",")) \
    .map(lambda cols: (cols[0], cols[1], cols[2],
float(cols[3]))) # assuming temp is in column 4

# 1. Filter for "TMIN" and find minimum temperatures
tmin_rdd = data_rdd.filter(lambda x: x[2] == "TMIN")

# a. Overall minimum temperature
overall_min_tmin = tmin_rdd.map(lambda x: x[3]).min()

# b. Minimum temperature for each ItemID
min_temp_by_item = tmin_rdd.map(lambda x: (x[1], x[3])) \
    .reduceByKey(lambda a, b: min(a, b))
```

```

# c. Minimum temperature for each StationID
min_temp_by_station = tmin_rdd.map(lambda x: (x[0], x[3])) \
                                .reduceByKey(lambda a, b: min(a, b))

# Display results for TMIN
print(f"Overall minimum temperature (TMIN): {overall_min_tmin}")
print("Minimum temperature for each ItemID (TMIN):")
for item, min_temp in min_temp_by_item.collect():
    print(f"ItemID: {item}, Min Temp: {min_temp}")

print("Minimum temperature for each StationID (TMIN):")
for station, min_temp in min_temp_by_station.collect():
    print(f"StationID: {station}, Min Temp: {min_temp}")

# 2. Filter for "TMAX" and find maximum temperatures
tmax_rdd = data_rdd.filter(lambda x: x[2] == "TMAX")

# a. Overall maximum temperature
overall_max_tmax = tmax_rdd.map(lambda x: x[3]).max()

# b. Maximum temperature for each ItemID
max_temp_by_item = tmax_rdd.map(lambda x: (x[1], x[3])) \
                             .reduceByKey(lambda a, b: max(a, b))

# c. Maximum temperature for each StationID
max_temp_by_station = tmax_rdd.map(lambda x: (x[0], x[3])) \
                               .reduceByKey(lambda a, b: max(a, b))

# Display results for TMAX
print(f"Overall maximum temperature (TMAX): {overall_max_tmax}")
print("Maximum temperature for each ItemID (TMAX):")
for item, max_temp in max_temp_by_item.collect():
    print(f"ItemID: {item}, Max Temp: {max_temp}")

print("Maximum temperature for each StationID (TMAX):")
for station, max_temp in max_temp_by_station.collect():
    print(f"StationID: {station}, Max Temp: {max_temp}")

```

Ex 7

Set up a simple Hadoop environment using Docker containers, including at least one NameNode and one DataNode. Ensure the containers are properly configured to interact with each other. After the setup, verify that the Hadoop cluster is operational by running a simple HDFS file operation (e.g., uploading a file to HDFS).

Check and install docker

```
docker --version
```

```
sudo apt install docker.io
```

```
sudo docker run hello-world
```

```
docker network create hadoop-net
```

```
docker pull bde2020/hadoop-namenode:2.0.0-hadoop3.2.1-java8
```

```
docker pull bde2020/hadoop-datanode:2.0.0-hadoop3.2.1-java8
```

```
docker run -d \
  --name namenode \
  --network hadoop-net \
  -e CLUSTER_NAME="my-hadoop-cluster" \
  -e CORE_CONF_fs_defaultFS=hdfs://namenode:9000 \
  bde2020/hadoop-namenode:2.0.0-hadoop3.2.1-java8
```

```
docker run -d \
  --name datanode \
  --network hadoop-net \
  -e CORE_CONF_fs_defaultFS=hdfs://namenode:9000 \
  bde2020/hadoop-datanode:2.0.0-hadoop3.2.1-java8
```

```
docker ps
```

```
docker logs namenode
```

```
docker logs datanode
```

(Run simple hdfs file system)

```
echo "Hello Hadoop!" > hello.txt
```

```
docker cp hello.txt namenode:/hello.txt
```

```
docker exec -it namenode hdfs dfs -mkdir -p /user/root  
docker exec -it namenode hdfs dfs -put /hello.txt /user/root/
```

```
docker exec -it namenode hdfs dfs -ls /user/root/
```

```
docker exec -it namenode hdfs dfs -cat /user/root/hello.txt
```

(clean up)

```
docker rm -f namenode datanode  
docker network rm hadoop-net
```

AWS

Creating Custom VPC, EC2 Instance and working on SG & NACL

1. Login into your AWS account.
2. Choose VPC Service
3. Choose the region Mumbai
4. Delete the existing VPC and setup custom VPC and its components
5. Get 2 elastic public IP
6. Create two EC2 instances and attach the public IP address
7. Name VM1 as Web Server & VM2 as Web Client
8. Connect to the instance via EC2 instance connect
9. Install Apache (web service) in Web Server
10. Install Links (web client) in Web Client
11. In the Security Group of Web Server, add rule to allow HTTP access.
12. Allow SSH & HTTP on the NACL
13. Test the web access from the web client using links app.

VPC and EC2

VPC : snu-dc-vpc
Ipv4 : 192.168.0.0/16

=> Connect subnet to vpc created

Subnet name: public-subnet

Availability zone: Mumbai

IPv4 subnet block (2nd one) : 192.168.1.0/24

=> Internet gateway

Name tag: snuc-dc-igw

=> Connect ig to vpc created

=> Route tables

Click on route table id -> edit routes -> add routes

0.0.0.0/0 - Internet gateway -> igw (auto-completion)

(chk point: completed setting up of vpc and it's components)

=> Elastic IP

Allocate ip address (create web server and web client (rename after creation))

=> EC2

Go to instances -> launch instances

Number of instances: 2

Name: VMs

Quick start -> change to ubuntu

Key pair -> proceed without key pair (not)

Now launch instance

-> go to instance

Rename to web server and web client

-> go to elastic ip

Click on web server ip add -> associate elastic ip add with it

Instance: choose web server

Do same for web client

(chk point: create two ec2 instances and attach public ip address)

-> go to instances

Click web server and connect -> connect instance

(if error ipv4 not public check again and refresh then connect)

=>New tab for terminal opens up(ec2-instance-connect)

```
ping 8.8.8.8
sudo apt update
sudo apt install apache2
service apache2 status (ctrl c to escape)
```

=> connect web client to instance
(tab should pop up)

```
ping 8.8.8.8
sudo apt update
sudo apt install links
clear
```

(chck point: Install apache and links)

Click on web server ip add -> go to security
Click security group id
Click on edit inbound rules

-> add rule -> http -> source -> anywhere ipv4

(chck point: HTTP access)

Go to instances -. Click link for web server go to security

=> nacl is in vpc

Vpc ->network security -> nacl

Click on nacl id
Click on edit inbound rules

Remove the existing rule
Add new rule

Type : SSH (20)

New rule -> 101 -> HTTP(80) -> save changes

(chck point: test the web access)

Go to ec2 interface for web client (the cleared one)

links (the public ip of web server -> in interface terminal)

```
ubuntu@ip-192-168-1-203:~$  
ubuntu@ip-192-168-1-203:~$  
ubuntu@ip-192-168-1-203:~$ links 13.126.225.0  
ubuntu@ip-192-168-1-203:~$  
ubuntu@ip-192-168-1-203:~$  
ubuntu@ip-192-168-1-203:~$ links http://13.126.225.0
```

Now removing the nacl from

its not working. let me check

Since Client and Server are in the same subnet. We cant get the NACL config. Lets delete the point

Removing the NACL config

Remove the created inbound rules (100,101)

Add new rule -> 100-> all traffic -> save changes

Now it will work

Its working.

If you have the Client and Server in different subnet you can configure NACL and test it. Becoz NACL is applied on Subnet Level.

(1. VPC & EC2 Lab : <https://youtu.be/AsSQb--MNXA> (no audio))

Route 53 Labs

Open route 53

Go to dashboard -> create hosted zone

=> inside hosted zone

Domain name:

Domain name [Info](#)
This is the name of the domain that you want to use.

Then create it

Go to go daddy domain

Log in

Now copy (.com) value/route traffic from route 53

-> add new record in godaddy

-> type: NS -> value (copied .com)

[NS records](#) determine which nameservers manage a domain's zone file.

| Type * | Name * | Value * |
|--------|-------------|-----------------------|
| NS | 21011101037 | ns-395.awsdns-49.com. |

Domains

Portfolio

DNS

Transfers

Services

Tools

Settings

Filters

Actions

| | | | | | | |
|--------------------------|----|-------------|-------------------------|-------------|--------------|------------|
| <input type="checkbox"/> | A | varun-todo | 3.108.220.63 | 600 seconds | | |
| <input type="checkbox"/> | NS | @ | ns13.domaincontrol.com. | 1 Hour | Can't delete | Can't edit |
| <input type="checkbox"/> | NS | @ | ns14.domaincontrol.com. | 1 Hour | Can't delete | Can't edit |
| <input type="checkbox"/> | NS | 21011101009 | ns-517.awsdns-00.net. | 1 Hour | | |
| <input type="checkbox"/> | NS | 21011101037 | ns-395.awsdns-49.com. | 1 Hour | | |
| <input type="checkbox"/> | NS | 21011101040 | ns-93.awsdns-11.com. | 1 Hour | | |

<<

<

1

2

3

4

...

>

>>

10

Put Your Domain to Work

Now go back

Go to hosted zones -> create records

Record name: www

Ip address: use ip of web server (public ipv4) for instances (created prev)

Create record

Go to web client terminal -> links url

URL : `www.<registration_no>.ngaws.xyz`

nslookup url

(<https://www.youtube.com/watch?v=-ndsfa-6GMI>)

IAM :

B.tech AI & Data Science

1.

a) Create IAM users(alice & bob), put them under group(server admin) and give full access to EC2 services

b) Create IAM users(cathy & david) put them under group(dns admin) and give full access to Route 53 services

c) Create IAM user(eve) give him access to billing

d) Create IAM usr(your_name) and give full access to all the services

e) Create an alias name for your account & check the login via alias url instead of account ID.

f) Login as alice & create an EC2 instance (Name : Web Server) with elastic public ip with required port numbers open in the security group.

g) check whether the port 80 is open for your instance from the below url

<https://portchecker.co/>

=> search for IAM

-> go to users



-> go to create user

alice -> provide user access click -> Specify user details

alice

The user name can have up to 64 characters. Valid characters: A-Z, a-z, 0-9, and * = , @ _ - (hyphen)

☒ Provide user access to the AWS Management Console - optional
If you're providing console access to a person, it's a [best practice](#) to manage their access in IAM Identity Center.

Are you providing console access to a person?

User type

☐ Specify a user in Identity Center - Recommended
We recommend that you use Identity Center to provide console access to a person. With Identity Center, you can centrally manage user access to their AWS accounts and cloud applications.

☒ I want to create an IAM user
We recommend that you create IAM users only if you need to enable programmatic access through access keys, service-specific credentials for AWS CodeCommit or Amazon Keyspaces, or a backup credential for emergency account access.

Console password

☒ Autogenerated password
You can view the password after you create the user.

☐ Custom password
Enter a custom password for the user.

☐ Show password

☒ Users must create a new password at next sign-in - Recommended
Users automatically get the `IAMUserChangePassword` policy to allow them to change their own password.

If you are creating programmatic access through access keys or service-specific credentials for AWS CodeCommit or Amazon Keyspaces, you can generate them after you create this IAM user. [Learn more](#)

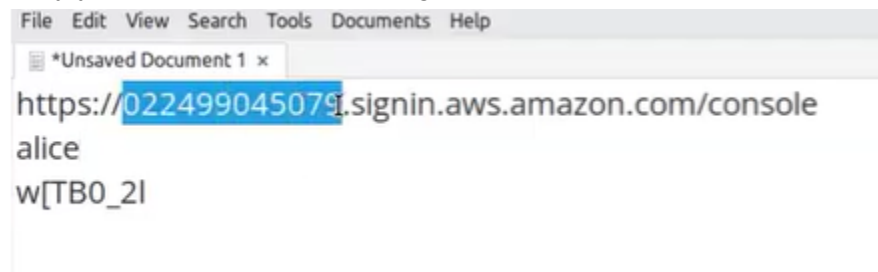
Cancel Next

Set permissions -> give next

Click on creator user after getting user details

Copy the password and Console sign-in URL

Copy your account id from top right (???)



Return to users list

Go to user groups

Now go to chrome and paste alice's url

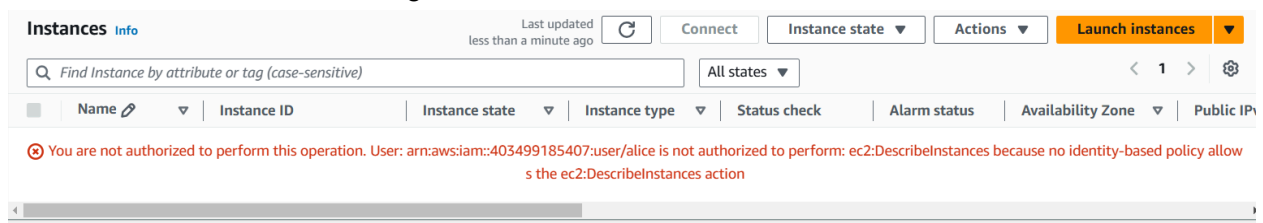
Now fill alice name and password (copied)

The image shows the AWS IAM console interface for changing a user's password. At the top left is the AWS logo. Below it, a message says "You must change your password to continue". The main form area has the following fields:
- AWS account: 403499185407
- IAM user name: alice
- Old password: [text input field]
- New password: [text input field]
- Retype new password: [text input field]
Below these fields is a blue button labeled "Confirm password change". At the bottom, there is a link that says "Sign in using root user email".

New password : old password+1


(put alice in a diff tab as you will be signed out of your actual account!)

=> search for ec2 on alice and go to instances



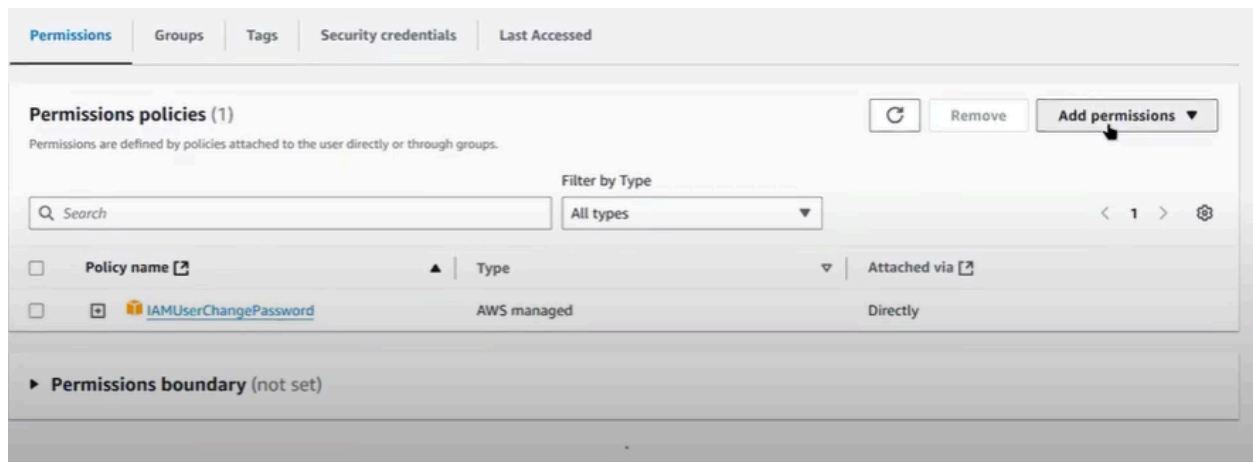
Now go to Cynddia => IAM

Go to policies




Go to users -> click on alice

Go to permissions -> add permissions x 2



Attach policies directly



Next -> add permission

Now alice will have access to it

Go to create user group (in cynddia)

Name the group

User group name
Enter a meaningful name to identify this group.

server-team

Maximum 128 characters. Use alphanumeric and '+', '@', '_' characters.

Add users to the group - Optional (1/1) info

An IAM user is an entity that you create in AWS to represent the person or application that uses it to interact with AWS.

Search

| <input checked="" type="checkbox"/> | User name | Groups | Last activity | Creation time |
|-------------------------------------|-----------|--------|----------------|----------------|
| <input checked="" type="checkbox"/> | alice | 0 | 11 minutes ago | 15 minutes ago |

Click alice and go for attached permission policies

Add users to the group - Optional (1/1) info

An IAM user is an entity that you create in AWS to represent the person or application that uses it to interact with AWS.

Search

| <input checked="" type="checkbox"/> | User name | Groups | Last activity | Creation time |
|-------------------------------------|-----------|--------|----------------|----------------|
| <input checked="" type="checkbox"/> | alice | 0 | 11 minutes ago | 15 minutes ago |

Attach permissions policies - Optional (1/955) info

You can attach up to 10 policies to this user group. All the users in this group will have permissions that are defined in the selected policies.

Filter by Type

ec2f All types 2 matches

| <input checked="" type="checkbox"/> | Policy name | Type | Used as | Description |
|-------------------------------------|-------------------------|-------------|---------|---|
| <input checked="" type="checkbox"/> | AmazonEC2FullAccess | AWS managed | None | Provides full access to Amazon EC2 via... |
| <input type="checkbox"/> | EC2FastLaunchFullAccess | AWS managed | None | This policy grants full access to EC2 Fa... |

Remove the policy

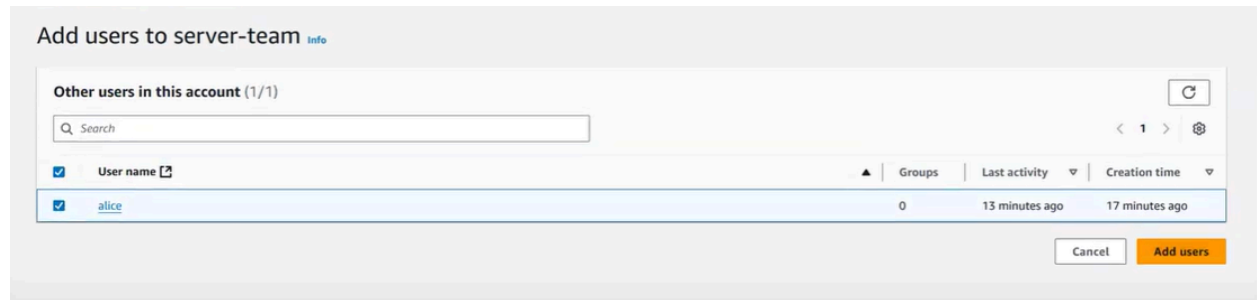
Remove permissions for user?

Remove user from group server-team?

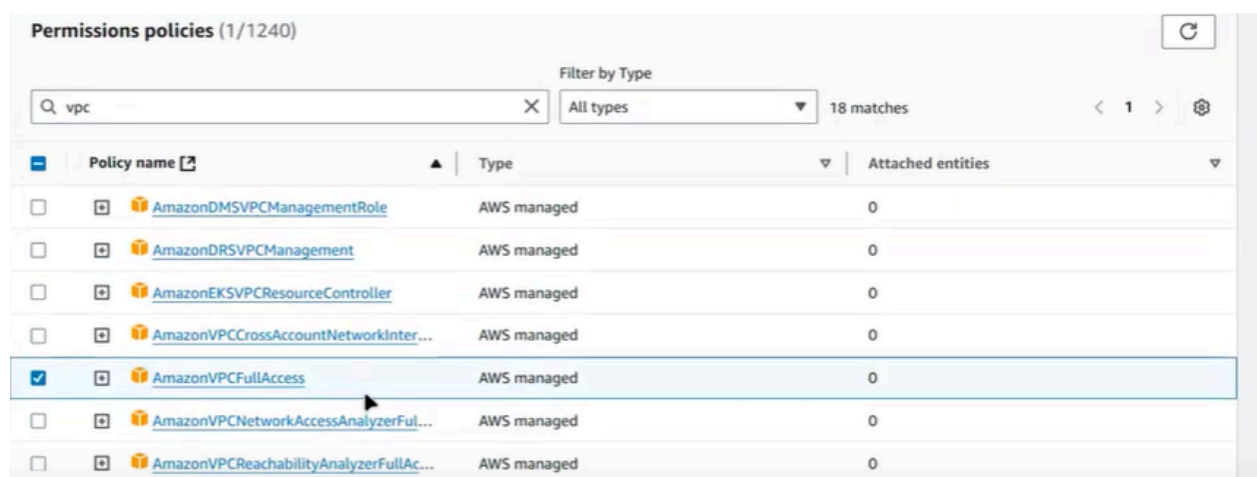
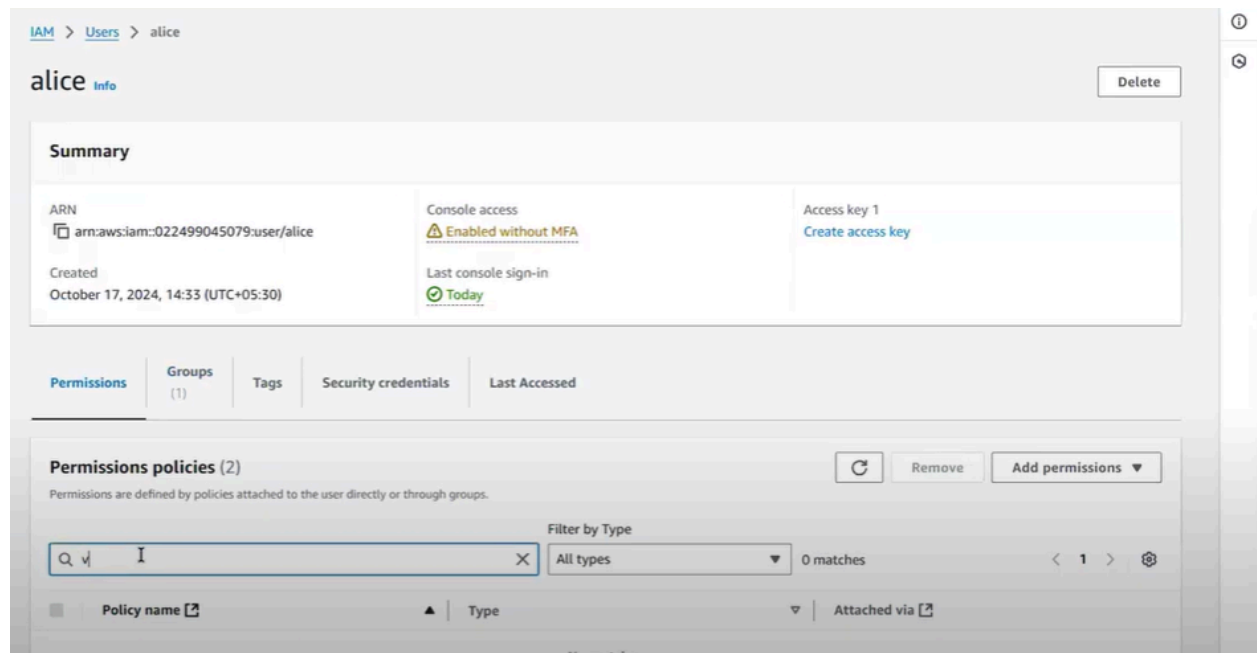
Remove policy AmazonEC2FullAccess?

Cancel Remove

Now click server teams in user groups



Add permission to alice in Cynddia



Create bob similarly

Search for vpc in bob

Create subnet

Subnet name
Create a tag with a key of 'Name' and a value that you specify.

test

The name can be up to 256 characters long.

Availability Zone [Info](#)
Choose the zone in which your subnet will reside, or let Amazon choose one for you.

No preference

IPv4 VPC CIDR block [Info](#)
Choose the VPC's IPv4 CIDR block for the subnet. The subnet's IPv4 CIDR must lie within this block.

192.168.0.0/16

IPv4 subnet CIDR block

192.168.10.0/24 256 IPs

< > ^ v

▼ **Tags - optional**

| Key | Value - optional | |
|--------|------------------|--------|
| Q Name | Q test | Remove |

Add new tag

You can add 49 more tags.

Remove

Add new subnet

Cancel Create subnet

Now to go your user

In user groups go for permission

And now go to route 53 in bob

Create host zone -> getting started

Services

Search

[Alt+S]

Domain name

Info

This is the name of the domain that you want to route traffic for.

testzone

Valid characters: a-z, 0-9, ! " # \$ % & ' () * + , - / : ; < = > ? @ [\] ^ _ ` { | } . ~

Description - optional

Info

This value lets you distinguish hosted zones that have the same name.

The hosted zone is used for...

The description can have up to 256 characters. 0/256

Type

Info

The type indicates whether you want to route traffic on the internet or in an Amazon VPC.

☒ Public hosted zone

A public hosted zone determines how traffic is routed on the internet.

☐ Private hosted zone

A private hosted zone determines how traffic is routed within an Amazon VPC.

Tags

Info

Apply tags to hosted zones to help organize and identify them.

No tags associated with the resource.

Add tag

You can add up to 50 more tags.

Cancel

Create hosted zone

Error , full user needs to give access

Permission in user groups

IAM > User groups > server-team

server-team [Info](#)

[Delete](#)

Summary

[Edit](#)

| | | |
|-----------------|-------------------------------------|---|
| User group name | Creation time | ARN |
| server-team | October 17, 2024, 14:48 (UTC+05:30) | arn:aws:iam::022499045079:group/server-team |

[Users](#) (2) | [Permissions](#) | [Last Accessed](#)

Users in this group (2)

An IAM user is an entity that you create in AWS to represent the person or application that uses it to interact with AWS.

[Refresh](#) [Remove](#) [Add users](#)

| <input type="checkbox"/> | User name | Groups | Last activity | Creation time |
|--------------------------|-----------------------|--------|----------------|----------------|
| <input type="checkbox"/> | alice | 1 | 24 minutes ago | 28 minutes ago |
| <input type="checkbox"/> | bob | 1 | 7 minutes ago | 8 minutes ago |

Attach permission policies to server-team

► **Current permissions policies (1)**

Other permission policies (1/954) [Refresh](#)

You can attach up to 10 managed policies to this user group. All of the users in this group inherit the attached permissions.

Filter by Type: [All types](#)

| <input type="checkbox"/> | Policy name | Type | Used as | Description |
|-------------------------------------|---|----------------------------|---------|--|
| <input checked="" type="checkbox"/> | AdministratorAccess | AWS managed - job function | None | Provides full access to AWS services an... |
| <input type="checkbox"/> | AdministratorAccess-Amplify | AWS managed | None | Grants account administrative permissi... |

Now create host zone for bob

Domain name [Info](#)

This is the name of the domain that you want to route traffic for.

Valid characters: a-z, 0-9, ! " # \$ % & ' () * + , - / : ; < = > ? @ [\] ^ _ ` { } . -

Description - optional [Info](#)

Create role in your user id

Identity and Access Management (IAM)

Search IAM

Dashboard

Access management

User groups

Users

Roles

Policies

Identity providers

Account settings

Access reports

Access Analyzer

External access

Unused access

Analyzer settings

Credential report

Organization activity

Service control policies

IAM > Roles

Roles (4) Info

Refresh

Delete

Create role

Search

< 1 >

Settings

| <input type="checkbox"/> | Role name | Trusted entities | Last activity |
|--------------------------|---|--------------------------------------|---------------|
| <input type="checkbox"/> | AWSServiceRoleForOrganizations | AWS Service: organizations (Service- | - |
| <input type="checkbox"/> | AWSServiceRoleForSSO | AWS Service: sso (Service-Linked Rol | 23 hours ago |
| <input type="checkbox"/> | AWSServiceRoleForSupport | AWS Service: support (Service-Link | - |
| <input type="checkbox"/> | AWSServiceRoleForTrustedAdvisor | AWS Service: trustedadvisor (Service | - |

Roles Anywhere Info

Manage

Authenticate your non AWS workloads and securely provide access to AWS services.

OS

OS

OS

Access AWS from your non AWS workloads

Operate your non AWS workloads using the same authentication and authorization strategy that you use within AWS.

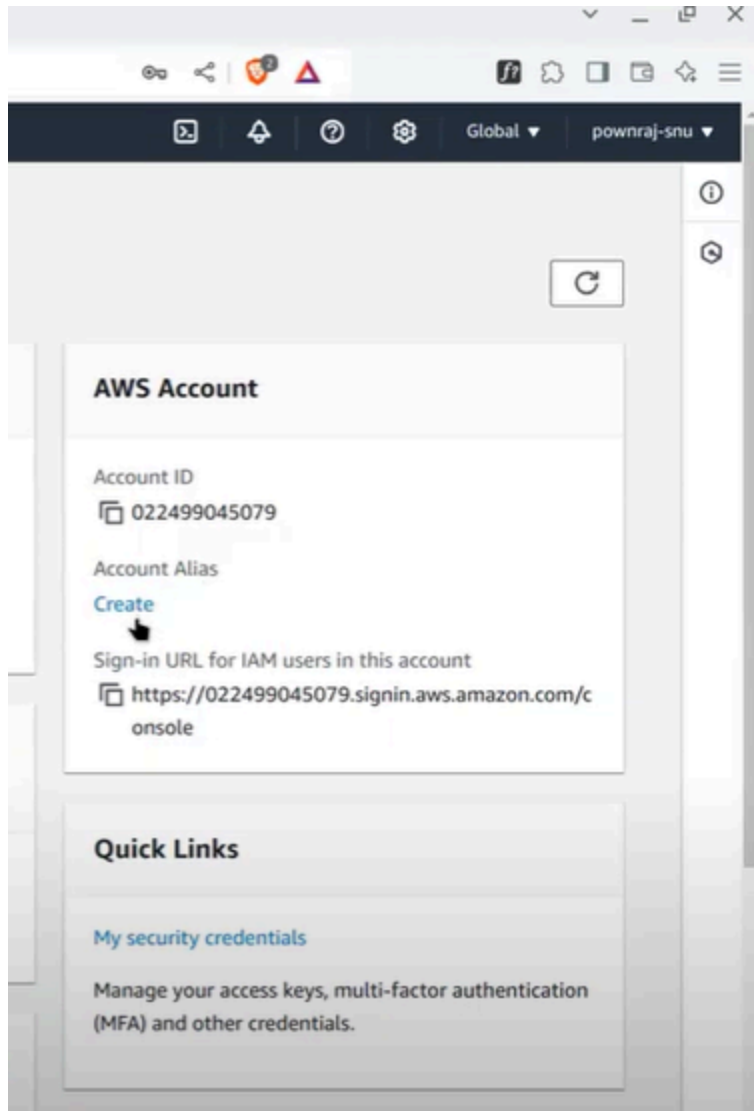
X.509 Standard

Use your own existing PKI infrastructure or use [AWS Certificate Manager Private Certificate Authority](#) to authenticate identities.

Temporary credentials

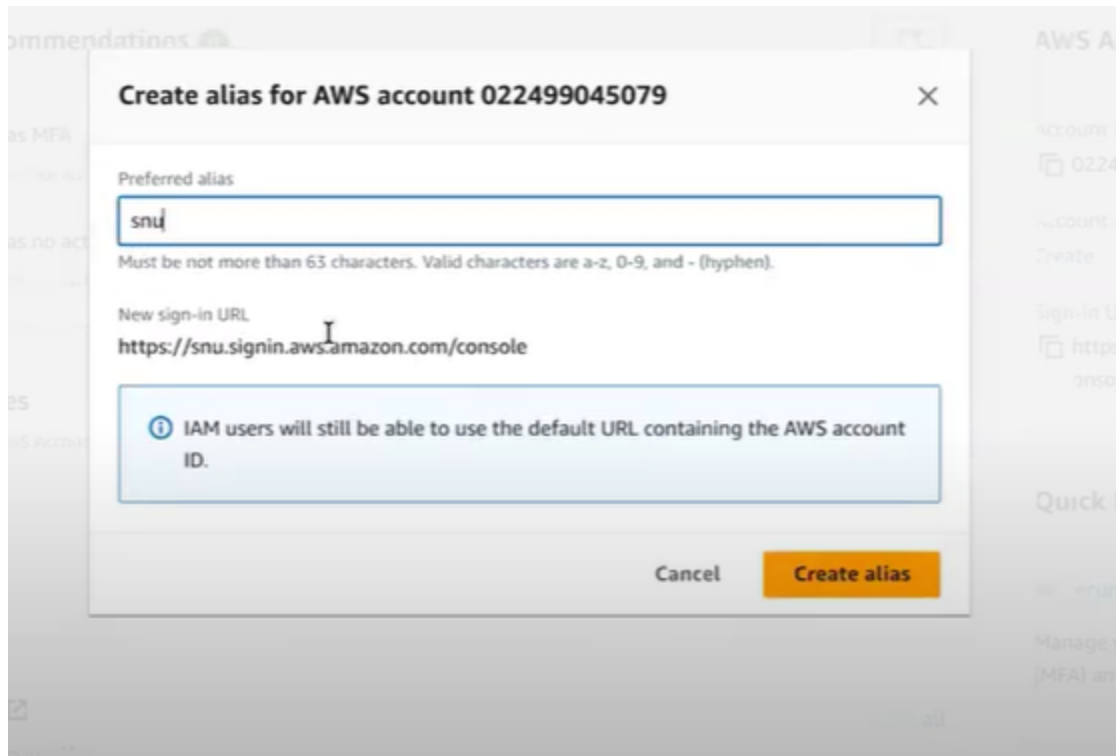
Use temporary credentials with ease and benefit from the enhanced security they provide.

Create alias for your account



Sign out of bob

Sign into alice



Instructions:

IAM :

Here's a structured guide to create IAM users and configure AWS access and services based on your requirements. I'll assume you have AWS administrative access to start this process:

Step-by-Step Guide

Step 1: Set Up IAM Groups and Users

1. **Log in to AWS Management Console:**
 - Go to the [IAM Console](#).
 2. **Create the Server Admin Group (Full EC2 Access):**
 - In the IAM Console, select **Groups** > **Create Group**.
 - Name it **ServerAdmin**.
 - Attach the **AmazonEC2FullAccess** policy.
 - Click **Create Group**.
 3. **Create Users Alice and Bob and Add to ServerAdmin Group:**
 - Go to **Users** > **Add Users**.
 - Add **Alice** and **Bob** as separate users.
 - Select **AWS Management Console Access** and configure their passwords.
 - Under **Permissions**, add **Alice** and **Bob** to the **ServerAdmin** group.
 - Click **Next**, review, and **Create Users**.
 4. **Create the DNS Admin Group (Full Route 53 Access):**
 - Go back to **Groups** > **Create Group**.
 - Name it **DNSAdmin**.
 - Attach the **AmazonRoute53FullAccess** policy.
 - Click **Create Group**.
 5. **Create Users Cathy and David and Add to DNSAdmin Group:**
 - Go to **Users** > **Add Users**.
 - Add **Cathy** and **David** as separate users.
 - Select **AWS Management Console Access** and set their passwords.
 - Under **Permissions**, add **Cathy** and **David** to the **DNSAdmin** group.
 - Review and **Create Users**.
 6. **Create the Billing User (Eve):**
 - Go to **Users** > **Add User**.
 - Name the user **Eve**.
 - Select **AWS Management Console Access**.
 - Assign the **AWSBillingReadOnlyAccess** policy directly to **Eve**.
 - Complete and **Create User**.
 7. **Create Your User with Full Access:**
 - Go to **Users** > **Add User**.
 - Name the user as **your_name** (replace with your actual name or preferred username).
 - Select **AWS Management Console Access**.
 - Attach the **AdministratorAccess** policy for full AWS access.
 - Review and **Create User**.
-

Step 2: Create Alias for Account Login URL

1. In **IAM Console**, go to **Account Settings** (you may need administrative permissions).

2. Under **Account Alias**, click **Create Account Alias**.
 3. Set an alias (e.g., `companyname-alias`).
 4. Save, and use the alias in your URL:
 - o Format: `https://<alias>.signin.aws.amazon.com/console`
 5. Use this URL to log in to verify that the alias works.
-

Step 3: Log In as Alice and Launch EC2 Instance

1. **Log in** to the AWS Console as **Alice**.
 2. Go to the **EC2 Dashboard**.
 3. Click **Launch Instance**:
 - o Name: `Web Server`
 - o Select an AMI, e.g., Amazon Linux 2.
 - o Choose an instance type (e.g., t2.micro for free-tier eligibility).
 - o Configure instance details as needed.
 - o **Assign Elastic IP**:
 - Under **Network settings**, select **Assign a public IP**.
 4. **Configure Security Group**:
 - o In the Security Group, ensure ports such as **80 (HTTP)** and **22 (SSH)** are open to the appropriate IP ranges (e.g., `0.0.0.0/0` for testing, though it's more secure to restrict IPs in a production environment).
 5. Review and **Launch** the instance.
-

Step 4: Verify Port 80 Accessibility

1. Once the instance is running, copy the **public IP** of the instance.
2. Visit [Port Checker](#).
3. Enter the public IP and **Port 80** to check if it is open.

This setup will help ensure that the IAM and EC2 configurations work as specified. Let me know if you run into any specific issues during the setup!

Rest of codes

<https://chatgpt.com/share/672901af-72b8-8013-8cf7-673d3607fbb1>

